

Evaluation Interview - Expert [*id:4*]

(fourth iteration)

SLEEC 1

Q1. Is the SLEEC rule correct and effective in preventing or reducing ARI's unethical behavior?

I think the base rule is overall correct. However, I'm not fully convinced about the last hedge clause. Let's review SLEEC 1a and 1b to gain a better understanding.

Q2. Are the principles correctly mapped to the SLEEC rule, the base rule, and the hedge clauses? Are there any principles not covered by this mapping?

In my opinion, the mapping of principles is correct.

Q3. Is the proposed rule correctly labeled as Social, Legal, Ethical, Empathetic, and Cultural?

Yes, I think the labeling works. But I would definitely add Social, since it involves interaction and communication with the user.

SLEEC 1a

Q1. Is the SLEEC rule correct and effective in preventing or reducing ARI's unethical behavior?

I think this rule makes sense, but I would remove 'It's time for exercising session' from the initial condition to avoid inconsistencies with SLEEC 1. The same applies for SLEEC 1b.

Q2. Are the principles correctly mapped to the SLEEC rule, the base rule, and the hedge clauses? Are there any principles not covered by this mapping?

Yes, I think the principles are well aligned here.

Q3. Is the proposed rule correctly labeled as Social, Legal, Ethical, Empathetic, and Cultural?

I would say yes.

SLEEC 1b

Q1. Is the SLEEC rule correct and effective in preventing or reducing ARI's unethical behavior?

In this case, I think the rule works logically, but it's a bit problematic in execution. If the user refuses to open the door while the room is too warm, keeping it closed could be dangerous. You should reflect on this.

Q2. Are the principles correctly mapped to the SLEEC rule, the base rule, and the hedge clauses? Are there any principles not covered by this mapping?

I think the principles mapping is correct.

Q3. Is the proposed rule correctly labeled as Social, Legal, Ethical, Empathetic, and Cultural?

In my opinion, the labels fit.

SLEEC 2

Q1. Is the SLEEC rule correct and effective in preventing or reducing ARI's unethical behavior?

The rule is mostly correct but linguistically ambiguous. The condition "the user is not exercising" could mean either that the user has finished the exercise or that they have stopped midway. Given that the rule applies during training time, the latter seems intended; the user has paused rather than completed the activity. This ambiguity could lead the robot to misinterpret the situation.

Moreover, the rule lacks an explicit mechanism to detect when the user is over-exercising or showing signs of physical strain. As it stands, the robot may continue encouraging or prompting the next exercise even when this could cause harm.

Q2. Are the principles correctly mapped to the SLEEC rule, the base rule, and the hedge clauses? Are there any principles not covered by this mapping?

The rule currently reflects the principles you identified.

Q3. Is the proposed rule correctly labeled as Social, Legal, Ethical, Empathetic, and Cultural?

Yes, the labeling is appropriate.

SLEEC 2a

Q1. Is the SLEEC rule correct and effective in preventing or reducing ARI's unethical behavior?

Yes, the rule is logically coherent and well-structured. However, the trigger phrase "complains about being tired" should be preferred over a generic "complains", as "complains" is too broad and ambiguous (it could refer to any type of discomfort or dissatisfaction). The proposed specialization clarifies that the trigger concerns fatigue rather than injury or pain.

Q2. Are the principles correctly mapped to the SLEEC rule, the base rule, and the hedge clauses? Are there any principles not covered by this mapping?

The ethical principles are correctly represented.

Q3. Is the proposed rule correctly labeled as Social, Legal, Ethical, Empathetic, and Cultural?

The rule clearly belongs to the Ethical and Empathetic categories.

SLEEC 3

Q1. Is the SLEEC rule correct and effective in preventing or reducing ARI's unethical behavior?

Yes, the rule is correct and conceptually clear.

A small improvement would be to specify what "user data" means, since it can include a wide range of information, vital signs, exercise performance, voice logs, historical records, etc. As suggested in the discussion, it might be better to use "user medical data" or "user historical data" to disambiguate, depending on the intended scope.

Q2. Are the principles correctly mapped to the SLEEC rule, the base rule, and the hedge clauses? Are there any principles not covered by this mapping?

Yes, the mapping is correct. However, I would add that Autonomy has it is respected because the system asks for consent from the user.

Q3. Is the proposed rule correctly labeled as Social, Legal, Ethical, Empathetic, and Cultural?

Yes, the rule belongs under Legal and Ethical.

SLEEC 4

Q1. Is the SLEEC rule correct and effective in preventing or reducing ARI's unethical behavior?

I think the rule is generally correct and can be effective, although I see a small issue with waking the user up. Even if the user is not in REM sleep, being awakened could still be unpleasant. In my opinion, the rule feels a bit too strict, almost like a form of torture if the robot wakes the person up abruptly just because it's meal time.

That said, if the goal is to support the user's diet and eating schedule, and especially if we are dealing with diabetic patients, then I think it's justified. In that specific context, ensuring the user eats at fixed times is important for their health, so it makes sense.

Q2. Are the principles correctly mapped to the SLEEC rule, the base rule, and the hedge clauses? Are there any principles not covered by this mapping?

Yes, I think the mapping of principles is correct.

Q3. Is the proposed rule correctly labeled as Social, Legal, Ethical, Empathetic, and Cultural?

In my opinion, the labeling is correct.

SLEEC 5

Q1. Is the SLEEC rule correct and effective in preventing or reducing ARI's unethical behavior?

I think the rule works well and is quite straightforward. It can happen at any time, so it's more general than the others, and I think that makes sense.

The only thing I noticed is that the part 'remind the user of meal time' and 'explain why the user cannot eat now' sound a bit repetitive. In my opinion, it would be enough to just say 'explain why you cannot eat at the moment.'

Q2. Are the principles correctly mapped to the SLEEC rule, the base rule, and the hedge clauses? Are there any principles not covered by this mapping?

Yes, I think the mapping is correct.

Q3. Is the proposed rule correctly labeled as Social, Legal, Ethical, Empathetic, and Cultural?

In my opinion, the labeling fits. However, I won't say it Social because of the explanation shows politeness and respect in communication, I would remove this label.

SLEEC 6

Q1. Is the SLEEC rule correct and effective in preventing or reducing ARI's unethical behavior?

I think this rule works well overall; it's basically the opposite of the previous one, and it makes sense in that context. Still, I believe it could include an additional condition, for example: if the user is particularly distressed, then allow them to eat anyway. That would make the rule more flexible and sensitive to emotional states.

Q2. Are the principles correctly mapped to the SLEEC rule, the base rule, and the hedge clauses? Are there any principles not covered by this mapping?

Yes, I think the mapping of principles is coherent.

Q3. Is the proposed rule correctly labeled as Social, Legal, Ethical, Empathetic, and Cultural?

Yes, I agree with the labeling.

Final questions

Q5. Are there any SLEEC rules you consider missing?

I do not identify any missing rules.

Anything to add?

Nothing to add.